


BST and Red-Black Trees(二叉搜索树和红黑树)

1.0 Overview

Binary Trees

- ▶ Recursive definition
 1. An empty tree is a binary tree
 2. A node with at most two child subtrees is a binary tree
 3. Only what you get from 1 by a finite number of applications of 2 is a binary tree

- ▶ Is this a binary tree?




data structures can support **dynamic set operations**

- Search
- Minimum
- Maximum
- Predecessor
- Successor
- Insert
- Delete

1.1 Binary Search Trees


BST Representation

- Represented by a linked data structure of nodes
- $root(T)$ points to the root of tree T
- Each node contains field:
 - key
 - $left$ - pointer to left child: root of left subtree
 - $right$ - pointer to right child : root of right subtree
 - p - pointer to parent. $p[root(T)] = NIL$ (optional)



Binary Search Tree Property

- Stored keys must satisfy the *binary search tree* property
1. $\forall y$ in left subtree of x ,
then $key[y] \leq key[x]$
 2. $\forall y$ in right subtree of x ,
then $key[y] \geq key[x]$




Inorder Traversal

- The *binary search tree* property allows the keys of a binary search tree to be printed, in (monotonically increasing) order, recursively

INORDERTREEWALK(x)

- 1: **if** $x \neq NIL$ **then**
- 2: INORDERTREEWALK($left[x]$)
- 3: print $key[x]$
- 4: INORDERTREEWALK($right[x]$)



- How long does the walk take?
- Can you prove its correctness?

Correctness of Inorder-Walk


- ▶ Must prove that it prints all elements, in order, and that it terminates
- ▶ 1. By **induction on size of tree**, **Size = 0**: Easy
- ▶ 2. **Size ≥ 1** :
 - a. Prints left subtree in order by induction
 - b. Prints root, which comes after all elements in left subtree (still in order)
 - c. Prints right subtree in order (all elements come after root, so still in order)

Preorder Traversal

- ▶ The **binary search tree** property allows the keys of a binary search tree to be printed recursively

PREORDERTREEWALK(x)

- 1: **if** $x \neq \text{NIL}$ **then**
- 2: print $\text{key}[x]$
- 3: PREORDERTREEWALK($\text{left}[x]$)
- 4: PREORDERTREEWALK($\text{right}[x]$)




- ▶ How long does the walk take?
- ▶ Can you prove its correctness?

Postorder Traversal

- ▶ The **binary search tree** property allows the keys of a binary search tree to be printed recursively

POSTORDERTREEWALK(x)

- 1: **if** $x \neq \text{NIL}$ **then**
- 2: POSTORDERTREEWALK($\text{left}[x]$)
- 3: POSTORDERTREEWALK($\text{right}[x]$)
- 4: print $\text{key}[x]$



- ▶ How long does the walk take?
- ▶ Can you prove its correctness?

Querying a Binary Search Tree


- ▶ All dynamic-set search operations can be supported in $O(h)$ time
- ▶ $h = \Theta(\lg n)$ for a **balanced** binary tree (and for an average tree built by adding nodes in random order.)
 1. Self-balanced binary search trees will have $h = \Theta(\lg n)$
 2. Examples of such trees, red-black tree, AVL tree, 2-3 tree
- ▶ $h = \Theta(n)$ for an unbalanced tree that resembles a linear chain of n nodes in the worst case

Tree Search

TREESearch(x, k),

x is a node, k is a value

- 1: **if** $x = \text{NIL}$ or $k = \text{key}[x]$ **then**
- 2: **return** x
- 3: **if** $k < \text{key}[x]$ **then**
- 4: **return** TREESearch($\text{left}[x], k$)
- 5: **else**
- 6: **return** TREESearch($\text{right}[x], k$)




- ▶ Running time: $O(h)$
- ▶ Aside: tail-recursion

Iterative Tree Search

ITERATIVE TREESearch(x, k)

- 1: **while** $x \neq \text{NIL}$ and $k \neq \text{key}[x]$ **do**
- 2: **if** $k < \text{key}[x]$ **then**
- 3: $x \leftarrow \text{left}[x]$
- 4: **else**
- 5: $x \leftarrow \text{right}[x]$
- 6: **return** x



- ▶ The **iterative** tree search is more **efficient** on most computers
- ▶ The **recursive** tree search is more **straightforward**

```

1 IterativeTreeSearch(x, k)
2 while x != NIL and k != key[x] do
3     if k < key[x] then
4         x <- left[x]
5     else
6         x <- right[x]
7 return x

```

Finding Min & Max

- ▶ The binary-search-tree property guarantees that:
 1. The **minimum** is located at the **left-most** node
 2. The **maximum** is located at the **right-most** node

TREEMINIMUM(*x*)

```

1: while left ≠ NIL do
2:   x ← left[x]
3: return x

```

TREEMAXIMUM(*x*)

```

1: while right ≠ NIL do
2:   x ← right[x]
3: return x


```

Predecessor and Successor

- ▶ Successor of node *x* is the **node *y* such that *key*[*y*] is the smallest key greater than *key*[*x*]**
- ▶ The successor of the largest key is NIL
- ▶ Search consists of two cases:
 1. **If node *x* has a non-empty right subtree**, then *x*'s successor is the minimum in the right subtree of *x*
 2. **If node *x* has an empty right subtree**, then:
 - a. As long as we move to the left up the tree (move up through right children), we are visiting smaller keys
 - b. *x*'s successor *y* is the node that *x* is the predecessor of (*x* is the maximum in *y*'s left subtree)
 - c. In other words, *x*'s successor *y*, is the lowest ancestor of *x* whose left child is also an ancestor of *x*

TREESUCCESSOR(x)


```
1: if  $right[x] \neq \text{NIL}$  then  
2:   return TREEMINIMUM( $right[x]$ )  
3:  $y \leftarrow p[x]$   
4: while  $y \neq \text{NIL}$  and  $x = right[y]$  do  
5:    $x \leftarrow y$   
6:    $y \leftarrow p[y]$   
7: return  $y$ 
```



- ▶ Code for **predecessor** is symmetric
- ▶ Running time: $O(h)$

Insertion and Deletion

- ▶ Change the dynamic set represented by a BST
- ▶ Ensure the **binary search tree** property holds after change
- ▶ Insertion is easier than deletion



BST insertion

```
1 TreeInsert(T,z)  
2 y <- NIL  
3 x <- root[T]  
4 while x != NIL do  
5   y <- x  
6   if key[z] < key[x] then  
7     x <- left[x]  
8   else  
9     x <- right[x]  
10 p[z] <- y  
11 if y = NIL then  
12   root[T] <- z  
13 else if key[z] < key[y]  
14 then  
15   left[y] <- z  
16 else  
17   right[y] <- z
```

- ▶ Initialization: $O(1)$
- ▶ While loop in lines 3-10 searches for place to insert z , maintaining parent y . This takes $O(h)$ time
- ▶ Lines 11-18 insert the value: $O(1)$
- ▶ Total: $O(h)$ time to insert a node

Sorting using BSTs

SORT(A)

- 1: **for** $i \leftarrow 1$ to n **do**
- 2: **TREEINSERT**($A[i]$)
- 3: **INORDERTREEWALK**($root$)

BST Deletion

- ▶ **Case 0:** if x has no children:
then remove x
- ▶ **Case 1:** if x has one child:
then make $p[x]$ point to child
- ▶ **Case 2:** if x has two children (subtrees):
then swap x with its successor
perform case 0 or case 1 to delete it
- ▶ Total: $O(h)$ time to delete a node

```

1  TreeDelete(T,z)
2  if left[z] = NIL or right[z] = NIL then
3      y <- z
4  else
5      y <- TreeSuccessor[z]
6  if left[y] != NIL then
7      x <- left[y]
8  else
9      x <- right[y]
10 if x != NIL then
11     p[x] <- p[y]
12 if p[y] = NIL then
13     root[T] <- x
14 else if y <- left[p[i]] then
15     left[p[y]] <- x
16 else
17     right[p[y]] <- x
18 if y != z then
19     key[z] <- key[y]
20 return y

```

- ▶ How do we know case 2 should go to case 0 or case 1 instead of back to case 2?
Because when x has 2 children, its successor is the minimum in its right subtree, and that successor has no left child (hence 0 or 1 child)
- ▶ Any other choice?
Equivalently, we could **swap with predecessor instead of successor**. It might be good to alternate to avoid creating lopsided tree


1.2 Red-Black Trees

Overview

- ▶ Red-black trees are a variation of binary search trees to ensure that the tree is **balanced**
- ▶ Height is $O(\lg n)$, where n is the number of nodes
- ▶ Operations take $O(\lg n)$ time in the **worst case**

Red-Black Tree

- Binary search tree + 1 bit per node: the attribute color, which is either red or black
- All other attributes of BSTs are inherited: key, left, right, and p
- All empty trees(leaves) are colored black
 - We use a single sentinel, nil, for all the leaves of red-black tree T , with $\text{color}[\text{nil}] = \text{black}$
 - The root's parent is also $\text{nil}[T]$




Red-Black Properties

- ▶ 1. Every node is either **red** or **black**
- ▶ 2. The **root** is **black**
- ▶ 3. Every **leaf** (*nil*) is **black**
 - Note: this means every real node has 2 children
- ▶ 4. If a node is **red**, then both its children are **black**
 - Note: can't have 2 consecutive **reds** on a path
- ▶ 5. For each node, all paths from the node to descendant leaves contain the same number of **black** nodes.

Height of a Red-Black Tree

- ▶ **Height of a node:**
Number of edges in a longest path to a leaf
- ▶ **Black-height of a node x , $bh(x)$:**
Number of black nodes (including $nil[T]$) on the path from x to leaf, not counting x
- ▶ Black-height of a red-black tree is the black-height of its root:
By Property 5, black height is well defined



- ▶ What is the minimum black-height of a node with height h ?
A height- h node has **black-height** $\geq h/2$
- ▶ Theorem: A red-black tree with n internal nodes has height **$h \leq 2 \lg(n + 1)$**
How do you suppose we'll prove this?

- ▶ Prove: n -node RB tree has height **$h \leq 2 \lg(n + 1)$**
- ▶ Claim: A subtree rooted at a node x contains at least **$2^{bh(x)} - 1$** internal nodes
- ▶ Proof by induction on height h

- ▶ Base step: x has height 0 (i.e., NULL leaf node)
 1. So $bh(x) = 0$
 2. So subtree contains $2^{bh(x)} - 1 = 2^0 - 1 = 0$ internal nodes (TRUE)
- ▶ Inductive step: x has positive height and 2 children
 1. Each child has black-height of $bh(x)$ or $bh(x) - 1$
 2. So the subtrees rooted at each child contain at least $2^{bh(x)-1} - 1$ internal nodes
 3. Thus subtree at x contains $(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ nodes (TRUE)

可以得到一个结论:


$$h \leq 2lg(n + 1)$$

RB Trees: Worst-Case Time

- ▶ So we've proved that a red-black tree has $O(lg\ n)$ height
- ▶ Corollary: These operations take $O(lg\ n)$ time:
 Minimum(), Maximum()
 Successor(), Predecessor()
 Search()
- ▶ Insert() and Delete():
 1. Will also take $O(lg\ n)$ time
 2. But will need special care since they modify tree

Rotation

- ▶ Our basic operation for changing tree structure is called **rotation**:



- ▶ Preserves BST key ordering
- ▶ $O(1)$ time...just changes some pointers

Insertion


► Insertion: the basic idea

1. Insert x into tree, **color x red**
2. **r-b property 2** could be violated (if x is root and red)
If so, no other property is violated, we make x black.
3. Otherwise, **r-b property 4** might be violated (if $p[x]$ red)
If so, move violation up tree until a place is found where it can be fixed
4. Total time will be $O(\lg n)$

Case1:

► Case 1: uncle is **red**:


In figures below, all Δ 's are **equal-black-height subtrees**



- **Change colors** of some nodes, preserving **r-b property 5**: all downward paths have equal **b.h.** The while loop now continues with x 's grandparent as the new x

Case2:


- Case 2: uncle is **black**
Node x is a **right child**



- Set $x = p[x]$. Transform to case 3 via a **left-rotation**
- This preserves property 5: all downward paths contain same number of black nodes

Case3:

- ▶ Case 3: uncle is **black**
Node x is a **left child**



- ▶ Perform some **color changes** and do a **right rotation**
- ▶ Again, preserves property 5: all downward paths contain same number of black nodes

Delete

Ordinary BST Delete


- ▶ Case 1: If vertex to be deleted is a **leaf**, just delete it
- ▶ Case 2: If vertex to be deleted has just **one child**, replace it with that child
- ▶ Case 3: If vertex to be deleted has **two children**, then swap it with its successor

Bottom-up Deletion

- ▶ Think of V as having **an extra unit of blackness**. This extra blackness must be absorbed into the tree (by a red node), or propagated up to the root and out of the tree
- ▶ There are **four cases** our examples and rules assume that V is a left child. There are symmetric cases for V as a right child
- ▶ The node just deleted was U
- ▶ The node that replaces it is V , which has an extra unit of blackness
- ▶ The parent of V is P
- ▶ The sibling of V is S

Case1:


- ▶ Case 1: V's sibling, S, is **red**



- ▶ NOT a terminal case One of the other cases will now apply
- ▶ All other cases apply when S is black

Case2:


- ▶ Case 2: V's sibling, S, is **black** and has **two black children**



- ▶ Recolor S to be red
- ▶ P absorbs V's extra blackness:
 1. If P is red, we're done
 2. If P is black, it now has extra blackness and problem has been propagated up the tree

Case3:


- Case 3: S is **black**, S's **right child is red**



- 1. Rotate S around P
- 2. Swap colors of S and P, and color S's right child black
- This is the terminal case we're done

Case4:

- Case 4: S is **black**, S's **right child is black** and S's **left child is red**



- 1. Rotate S's left child around S
- 2. Swap color of S and S's left child before rotation
- 3. Now in case 3. e.g., V's sibling is black, which has a red right child.

1.3 Augmenting Data Structures

Overview

- ▶ A "textbook" data structure is enough in some situations
- ▶ Many others require a dash of creativity and it will suffice to **augment** a textbook data structure by storing **additional information** in it
- ▶ The added information must be **updated and maintained by the ordinary operations** on the data structure
- ▶ Then we can program **new operations** for the data structure to support the desired application

Dynamic Order Statistics

- ▶ We want to augment red-black trees so that they can support fast **order-statistic** operations
- ▶ So introducing an augmenting data structure: order-statistic tree:
 1. Besides the usual red-black tree fields $key[x]$, $color[x]$, $p[x]$, $left[x]$, and $right[x]$ in a node x , it has additional information: field **$size[x]$**
 2. $Size[x]$ is the number of (internal) nodes in the subtree rooted at x (including x itself)
 3. $Size[x] = size[left[x]] + size[right[x]] + 1$ ($size[nil[T]] = 0$)

- ▶ Keys need not to be distinct in an order-statistic tree
- ▶ In the presence of equal keys, it is well defined that the rank of an element is the position at which it would be printed in an inorder walk of the tree

