

# Lecture 1.9: 高斯混合模型与 EM 算法

2025.12.12

*Lecturer:* 宋骥

*Scribe:*

## 1 Gaussian Mixture Model

### 1.1 Mixture Model

MM is a probabilistic model that contains  $k$  sub-distribution among the overall distribution.

The  $k$  sub-model is the Hidden Variable of the Mixture model.

A MM can use any distribution.

GMM: a MM that uses Gaussian distribution as the sub-model.

Example:

$X_i$  - height of all students.

$z_i \in \{1, 2\}$  1: Boys, 2: Girls.

Distribution of  $X_i$  is not a simple normal (Gaussian) distribution, but a “**Mixture**” of two Gaussian distributions.

$$X_1 \sim N(\mu_1, \sigma_1^2), X_2 \sim N(\mu_2, \sigma_2^2)$$

$$X \sim (1 - \Delta) \cdot X_1 + \Delta \cdot X_2$$

$$\Delta \in \{0, 1\}, Pr(\Delta = 1) = \alpha$$

Define  $\phi_{\theta_i}(X)$  as the normal density (正态密度)

the probabilistic distribution of GMM is

$$P(X|\theta) = (1 - \alpha)\phi_{\theta_1}(X) + \alpha\phi_{\theta_2}(X)$$

$$\theta = (\alpha, \theta_1, \theta_2) = (\alpha, \mu_1, \sigma_1^2, \mu_2, \sigma_2^2)$$

Similarly, for k-GMM

$$P(X|\theta) = \sum_i^k \alpha_i \phi(X|\theta_i)$$

$$\theta = (\alpha, \mu, \sigma^2) = (\alpha_1, \dots, \alpha_k, \mu_1, \dots, \mu_k, \sigma_1^2, \dots, \sigma_k^2)$$

## 2 Estimating GMM parameters

### 2.1 Maximum Likelihood Estimation for Gaussian Model

Let's start with estimating the parameters of a single GM.

极大似然估计：

- 求似然函数，取对数.
- 对对数求导并使之为 0，解对数似然方程.  
即为未知参数的最大似然估计值.

For a Gaussian distribution  $\mathcal{X} \sim N(\mu, \sigma^2)$

Given  $n$  observable samples.

$$\begin{aligned} L(\mu, \sigma^2) &= \prod_{i=1}^n f(x_i; \mu, \sigma^2) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \exp\left(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}\right) \\ &= (2\pi\sigma^2)^{-\frac{n}{2}} \exp\left(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}\right) \end{aligned}$$

$$\begin{aligned} \ln L(\mu, \sigma^2) &= \ln \left[ (2\pi\sigma^2)^{-\frac{n}{2}} \exp\left(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}\right) \right] \\ &= \ln [(2\pi\sigma^2)^{-\frac{n}{2}}] + \ln \left[ \exp\left(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}\right) \right] \\ &= -\frac{n}{2} \ln(2\pi\sigma^2) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2} \\ &= -\frac{n}{2} \ln(2\pi) - \frac{n}{2} \ln(\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \end{aligned}$$

$$\begin{cases} \frac{\partial \ln L}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu) = 0 \\ \frac{\partial \ln L}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \mu)^2 = 0 \end{cases}$$

$$\Downarrow$$

$$\begin{cases} \sum_{i=1}^n (x_i - \mu) = 0 \\ \sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2 \end{cases}$$

$$\therefore \mu = \frac{1}{n} \sum_{i=1}^n x_i = \bar{x}$$

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

## 2.2 MLE for GMM

For GMM, we don't know a certain data sample belongs to which sub-model. We need to estimate  $\alpha, \mu, \sigma^2$ .

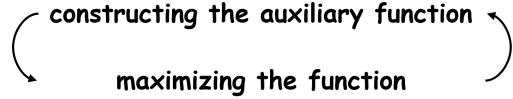
$$\begin{aligned} \log L(\theta) &= \sum_{j=1}^N \log[P(X_j|\theta)] \\ &= \sum_{j=1}^N \log\left[\sum_i^k \alpha_i \phi(X|\theta_i)\right] \end{aligned}$$

$\log \sum$  is not fun, there is no way to reduce it !

Expectation-Maximization (EM) procedure can be used to handle  $\log \sum$ .

① Summarization of EM: It uses Jensen's inequality to create a lower bound (called an auxiliary function) for the likelihood that uses  $\sum \log$  instead.

Repeat:



leading to a local maximum of the log likelihood for GMM.

Let's start with the overall process: assign each data point to a cluster, for each cluster, consider it as a Gaussian Model, calculate the  $\theta$ .

- First, generate the cluster assignment.

$$z_i|w \sim \text{Categorical}(w)$$

$w_k$  is the probability that  $i$ 's cluster is  $k$ .

$$P(z_i = k|w) = w_k$$

We call  $w_k$  the mixture weights

$$\sum_k w_k = 1, 0 \leq w_k \leq 1$$

- Then, generate  $x_i$  from the cluster's distribution:

$$x_i|z_i = k \sim N(\mu_k, \sum_k)$$

$$\begin{aligned} \log likelihood(\theta) &= \log \prod_i \sum_k P(X_i = x_i | z_i = k, \theta) P(z_i = k | \theta) \\ &= \sum_i \log \sum_k P(X_i = x_i | z_i = k, \theta) P(z_i = k | \theta) \end{aligned}$$

$z_i$  is the hidden variable, taking values  $i = 1, \dots, k$ .

Similarly:

$$\begin{aligned} \log likelihood(\theta) &= \sum_i \log \sum_k P(X_i = x_i, z_i = k | \theta) \\ &= \sum_i \log \sum_k P(z_i = k | \theta) P(X_i = x_i | z_i = k, \theta) \\ &= \sum_i \log \sum_k P(z_i = k | x_i, \theta_t) \frac{P(X_i = x_i, z_i = k | \theta)}{P(z_i = k | x_i, \theta_t)} \\ &= \sum_i \log E_z \frac{P(X_i = x_i, z_i = k | \theta)}{P(z_i = k | x_i, \theta_t)} \end{aligned}$$

**Lemma(Jensen's Inequality).** If  $f$  is convex, then  $f(EX) \leq E(f(X))$ . If  $f$  is convex,  $-f$  is concave, thus  $-f(EX) \geq -E(f(X)) = E(-f(X))$ .

Here,  $-f(X) = \log(X)$  which is concave, thus  $\log(EX) \geq E \log X$ .

$$\begin{aligned} \log likelihood(\theta) &= \sum_i \log E_z \frac{P(X_i = x_i, z_i = k | \theta)}{P(z_i = k | x_i, \theta_t)} \\ &\geq \sum_i E_z \log \frac{P(X_i = x_i, z_i = k | \theta)}{P(z_i = k | x_i, \theta_t)} \quad (\text{Jensen's Inequality}) \\ &= \sum_t \sum_k P(z_i = k | x_i, \theta_t) \log \frac{P(X_i = x_i, z_i = k | \theta)}{P(z_i = k | x_i, \theta_t)} \\ &=: A(\theta, \theta_t) \end{aligned}$$

- E-step: compute  $P(z_i = k|x_i, \theta_t) =: \gamma_{ik}$  for each  $i, k$
- M-step:  $\max_{\theta} A(\theta, \theta_t) = \sum_i \sum_j \gamma_{ik} \log \frac{P(X_i=x_i, z_i=k|\theta)}{\gamma_{ik}}$ ,  $\gamma_{ik}$  doesn't depend on  $\theta$ ,  
thus,  $\max_{\theta} \sum_i \sum_j \gamma_{ik} \log P(X_i = x_i, z_i = k|\theta)$ .

## ② Back to GMM

- E-step: Using Bayes Rule

$$\begin{aligned} P(z_i = k|x_i, \theta_t) &= \frac{P(X_i=x_i|z_i=k, \theta_t)P(z_i=k|\theta_k)}{P(X_i=x_i|\theta_t)} \\ &= \frac{N(x_i; \mu_{kt}, \sum_{k't})w_{kt}}{\sum_{k'} N(x_i; \mu_{k't}, \sum_{k't})w_{k't}} =: \gamma_{ik} \end{aligned}$$

- M-step:  $\max_{\theta} A(\theta, \theta_t) = \sum_i \sum_j \gamma_{ik} \log P(X_i = x_i, z_i = k|\theta)$

Update  $\theta$ , which is the collection  $w, \mu, \sum$ , by setting derivative of  $A$  to 0 with constraint  $\sum_k w_k = 1$

setting the derivatives to 0, we get the result for the cluster means:

$$\mu_{k,t+1} = \frac{\sum_i x_i \gamma_{ik}}{\sum_i \gamma_{ik}}$$

$$\text{Similarly, } \sum_{k,t+1} = \frac{\sum_i \gamma_{ik} (x_i - \mu_{k,t+1})(x_i - \mu_{k,t+1})^T}{\sum_i \gamma_{ik}}$$

Update for  $w$  is trickier because of the constraint. We need to use Lagrangian to solve this constrained optimization problem.

The Lagrangian:

$$L(\theta, \theta_t) = A(\theta, \theta_t) + \lambda(1 - \sum_k w_k)$$

Taking the derivative

$$\begin{aligned} \frac{\partial L(\theta, \theta_t)}{\partial w_{k'}} &= \frac{\partial A(\theta, \theta_t)}{\partial w_{k'}} - \lambda \\ &= \frac{\partial}{\partial w_{k'}} (\sum_i \sum_k \gamma_{ik} \log P(X_i = x_i, Z_i = k|\theta)) - \lambda \end{aligned} \tag{1}$$

By the probabilistic model for generating data according to GMM:

$$\begin{aligned} P(X_i = x_i, Z_i = k|\theta) &= P(Z_i = k|w) \cdot P(X_i = x_i|Z_i = k, \mu_{k,t+1}, \sum_{k,t+1}) - \lambda \\ &= w_k \cdot N(x_i; \mu_{k,t+1}, \sum_{k,t+1}) \end{aligned}$$

plugging back to (1)

$$\begin{aligned}\frac{\partial L(\theta, \theta_t)}{\partial w_{k'}} &= \sum_i \frac{\partial}{\partial w_{k'}} [\gamma_{ik'} \log[w_{k'} N(X; \mu_{k',t+1}, \sum_{k',t+1})]] - \lambda \\ &= \sum_i \frac{\partial}{\partial w_{k'}} [\gamma_{ik'} \log(w_{k',t+1})] + \frac{\partial}{\partial w_{k'}} [N(X, \mu_{k',t+1}, \sum_{k',t+1})]\end{aligned}$$

Here,  $N(X, \mu_{k',t+1}, \sum_{k',t+1})$  dose not depend on  $w_{k'}$

$$\begin{aligned}\frac{\partial L(\theta, \theta_t)}{\partial w_{k'}} &= \sum_i \frac{\partial}{\partial w_{k'}} [\gamma_{ik'} \log(w_{k'})] - \lambda \\ &= \sum_i \gamma_{ik'} \frac{1}{w_{k'}} - \lambda \\ &= \frac{1}{w_{k'}} \sum_i \gamma_{ik'} - \lambda\end{aligned}$$

Setting the derivative to 0,

$$w_{k',t+1} = \frac{\sum_i \gamma_{ik'}}{\lambda}$$

We know  $\sum_{k'} w_{k',t+1} = 1$ , so  $\lambda$  is the normalization factor

$$\begin{aligned}\lambda &= \sum_k \sum_i \gamma_{ik} = \sum_i \underbrace{\left( \sum_k P(Z_i = k | x_i; \theta) \right)}_{=1 \text{ as it is the sum over the whole probability distribution.}} \\ &= \sum_i 1 \\ &= n\end{aligned}$$

Finally  $w_{k',t+1} = \frac{\sum_i \gamma_{ik'}}{n}$

To sum:

- E-step : What is the current estimate of the probability that  $x_i$  comes from duster  $k$  ?  
 $\rightarrow \gamma_{ik}$
- M-step : update  $\mu, \sum, w$ .

### 3 Multinomial Mixture Models and its estimation(多项式混合模型)

EM can be used for GMM, similar idea can be used for k-means.

Task: Document review for a law firm, we need to scan tens of thousands of emails to determine their relevance to a legal matter.

It could be helpful to pre-group them according to their topics.

Multinomial Mixture Model(MMM) would model each topic as a distribution over English words.

### I. MMM

Let  $x_1, x_2, \dots, x_N$  be per-document word count vectors, where  $n$  is the total number of documents.

$X_{iv} = m$  if  $v$  appears  $m$  times in document  $i$ .

We assume each document contains  $M$  words in total, i.e.  $\sum_v X_{iv} = M$ .

The MMM models all documents as independent draws from the probability mass function

$$P(X) = \sum_{j=1}^k \pi_j f(x_j; \theta_j)$$

Each component

$$f(x_j; \theta_j) \propto \prod_{v=1}^p \theta_{jv}^{x_v}$$

is the pmf of a multinomial distribution,  $Mult(\theta_j, M)$ , where  $\theta_j \in \mathbb{R}^p$

$p$  is the total number of possible words,

$\theta_{jv}$  is the unknown probability of selecting word  $v$  for topic  $j$ .

### II. EM for MMMs.

$$\begin{aligned} \log p(X_{1:n}, Z_{1:n}; \pi, \theta_{1:k}) &= \sum_{i=1}^n \log p(X_i, Z_i; \pi, \theta_{1:k}) \\ &= \sum_{i=1}^n \log p(Z_i; \pi) + \sum_{i=1}^n \log p(X_i | Z_i; \theta_{i:k}) \\ &= \sum_{i=1}^n \log \pi_{Z_i} + \sum_{i=1}^n \sum_{v=1}^p x_{iv} \log \theta_{Z_i v} + \text{parameter-free terms} \\ &= \sum_{i=1}^n \sum_{j=1}^k \prod_{Z_{i:j}} \log \pi_j + \sum_{i=1}^n \sum_{j=1}^k \sum_{v=1}^p x_{iv} \prod_{Z_{i:j}} \log \theta_{jv} + \dots \end{aligned}$$

$\pi_j$  is the unknown probability of selecting topic  $j$ ,  $\sum_{j=1}^k \pi_j = 1$

· E-step

$$\mathbb{E}_{p(\cdot | X_{1:n}, \pi, \theta)} [\log p(X_{1:n}, Z_{1:n}; \pi_j, \theta_{1:k})] = \sum_{i=1}^n \sum_{j=1}^k \tau_{ij} \log \pi_j + \sum_{i=1}^n \sum_{j=1}^k \sum_{v=1}^p x_{iv} \tau_{ij} \log \theta_{jv}$$

where  $\tau_{ij} = p(Z_i = j \mid X_i; \pi, \theta_{1:k})$

By Bayes's Rule:

$$p(Z_i = j \mid X_i; \pi, \theta_{1:k}) \propto p(Z_i = j; \pi) \cdot p(X_i \mid Z_i, \theta_{1:k})$$

$$\therefore p(Z_i = j \mid X_i, \pi, \theta_{1:k}) = \frac{\pi_j f(X_i; \theta_j)}{\sum_{l=1}^k \pi_l f(X_i; \theta_l)}$$

• The M-step:

$$\pi_j = \frac{1}{n} \sum_{i=1}^n \tau_{ij}$$

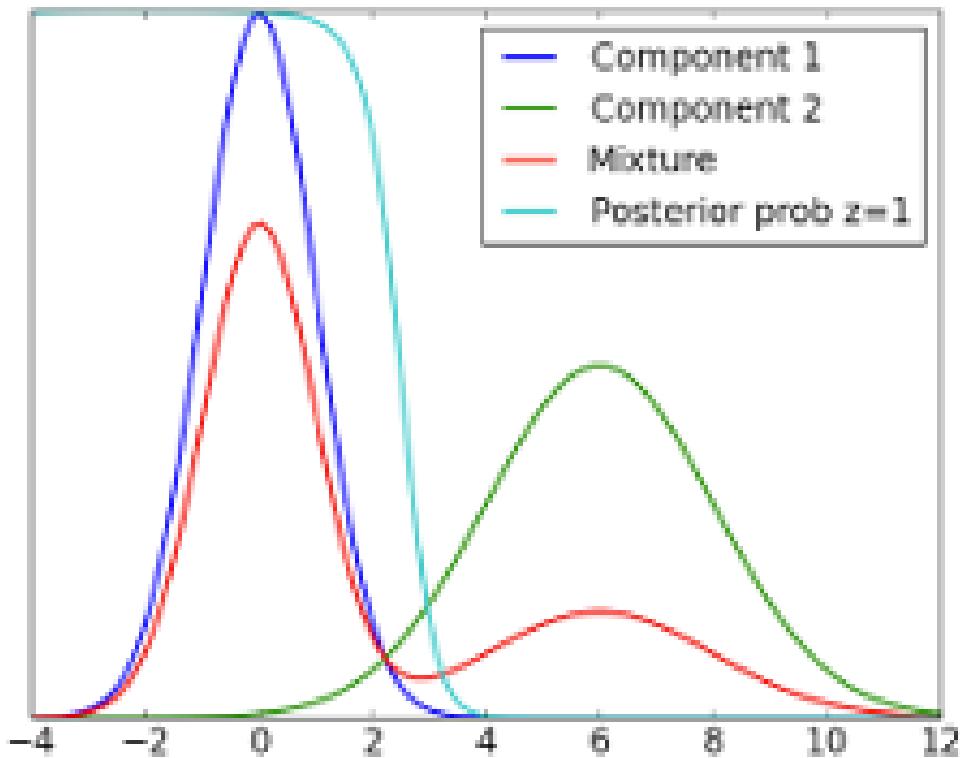
follow the similar steps of EM for GMM.

## 4 Model Inference \* 后验概率推理

Assume we have the parameters for each sub-model. Given a data point, determine which sub-model it belongs to.

$$p(z|x) \propto p(z)p(x|z)$$

Example 1:



$x = 2$ , with two GMs,

$$Pr(z = 1) = 0.7$$

$$Pr(z = 2) = 0.3$$

$$p(x|z = 1) = Guassian(2; 0.1) \approx 0.054$$

$$p(x|z = 2) = Guassian(2; 6.2) \approx 0.027$$

$$p(z = 1|x) = \frac{Pr(z=1)p(x|z=1)}{Pr(z=1)p(x|z=1)+Pr(z=2)p(x|z=2)} = 0.824$$

## 5 Conclusion

Practical considerations for EM:

- GMMs are not appropriate for all data types (e.g., discrete data).

We can use other models as the mixture model.

- The solutions of the EM algorithm are usually suboptimal.

Initialization strategies can help.

- No single satisfying approach to choosing  $k$ .

Methods we developed for  $k$ -means also apply to the GMM.

- We assume all data points are independent drawn from a common mixture.

What if not?  $\Rightarrow$  Hidden Markov models.